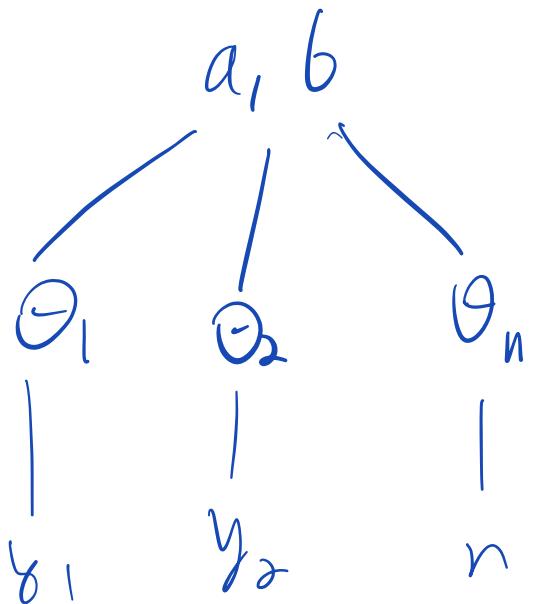


Lecture 5: Monte Carlo Methods

Professor Alexander Franks

Announcements

- Reading: Chapter 10 and 11
- HW 3 posted.



$$y_i \sim \text{Bin}(n_i, \theta_i)$$

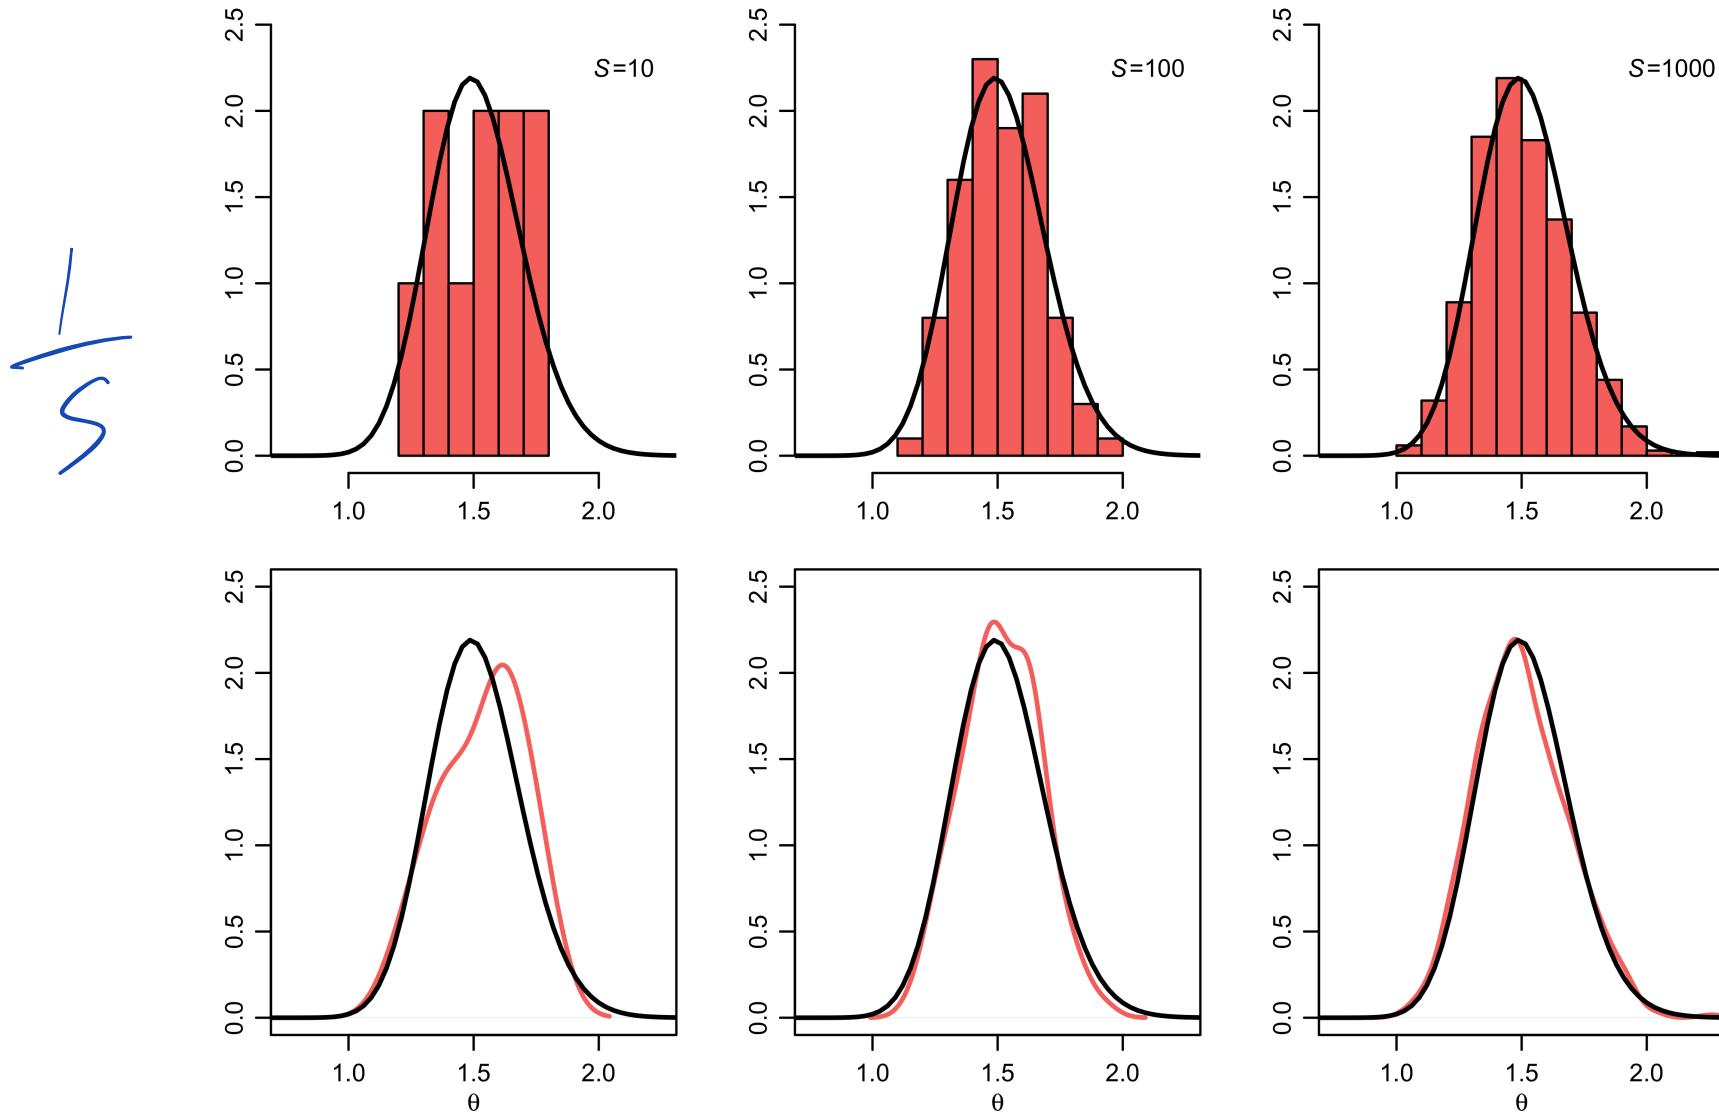
$$\theta_i \sim \text{Beta}(a, b)$$

Want: $P(a, b | y_1, \dots, y_n) =$

$$\frac{\int \int \int P(a, b, \theta_1, \dots, \theta_n | y_1, \dots, y_n) ds}{\alpha}$$

$$\frac{P(y_1, \dots, y_n | \theta_1, \dots, \theta_n, a, b) P(a, b, \theta^*)}{\left(\prod_{i=1}^{50} P(y_i | \theta_i) \right) P(\theta_1, \dots, \theta_{50} | a, b) P(a, b)}$$

Monte Carlo approximations of a distribution



Sampling strategies

- Inversion Sampling (works for univariate distns)
- Grid sampling (works for low dimensional problems)
- Rejection sampling
- Importance sampling
- Markov Chain Monte Carlo

Sampling strategies

Why sampling?

Probability Integral Transform

- Suppose that a random variable, Y has a continuous distribution for with CDF is F_Y .
- Then the random variable $U = F_Y(Y)$ has a uniform distribution
 - This is known as the “probability integral transform PIT”
- By taking the inverse of F_Y we have $F^{-1}(U) = Y$

Inversion Sampling

The inverse transform sampling method works as follows:

1. Generate a random number u from $\text{Unif}[0, 1]$
2. Find the inverse of the desired CDF, e.g. $F_Y^{-1}(u)$.
3. Compute $y = F_Y^{-1}(u)$. y is now a sample from the desired distribution.

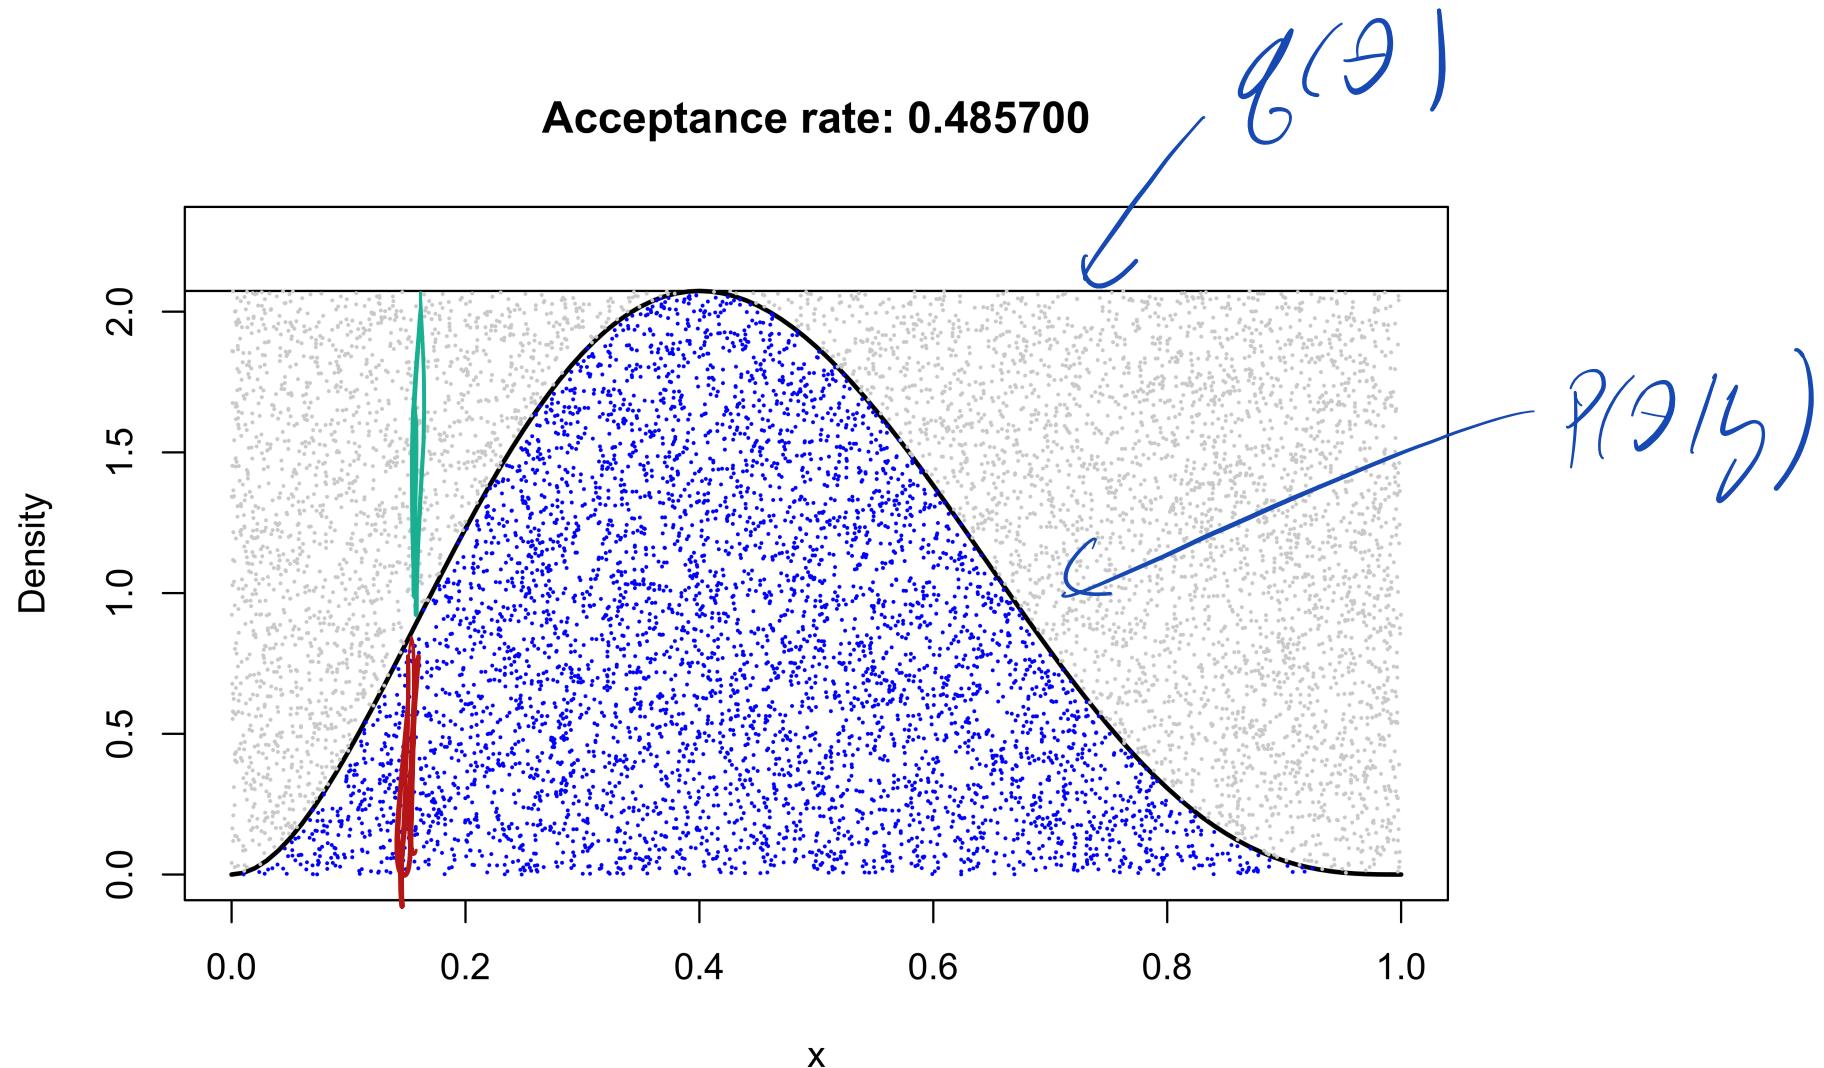
Inversion Sampling

Animation

Inversion Sampling

- Inversion sampling can be a fast and simple way to sample from a distribution
- Only effective if we know the inverse-CDF and can easily compute it
- This is a big challenge in practice. For example, even the normal distribution has a CDF, Φ , which cannot be expressed analytically.
 - Shifts from one hard problem (sampling) to another (computing an integral)
 - Need alternatives!

Rejection Sampling



Rejection Sampling algorithm

1. Choose a proposal density, $q(\theta)$ that we can easily sample from (e.g. uniform or normal) such that:

$$2. \text{Find } M = \max \frac{p(\theta|y)}{q(\theta)}$$

- If $M = \infty$ then q cannot be used as a proposal distribution
- If M is finite, $Mq(\theta)$ “envelopes” $p(\theta|y)$

3. Draw a sample, $\theta^{(s)}$ from $q(\theta)$

4. Accept $\theta^{(s)}$ as a draw from $p(\theta | y)$ with probability $\frac{p(\theta^{(s)}|y)}{Mq(\theta^{(s)})}$

Rejection Sampling

Demo

Rejection Sampling

- What is the expected number of rejections?
- How can we show this algorithm is valid?

Z be a random indicator for accepting
 θ^* proposed from g targeting P .

$$P(Z=1) = E_g[Z] = E[E[Z|\theta^*]]$$

$$= E_g \left[\frac{P(\theta^*|y)}{Mg(\theta^*)} \right] = \int \frac{P(\theta|y)}{Mg(\theta^*)} g(\theta^*) d\theta^*$$

$$= 1/M$$

$$\tilde{P}(\theta|y) = L(\theta)P(\theta)$$

$$\tilde{M} = \max \frac{\tilde{P}}{q} \quad E_g \left[\frac{\tilde{P}}{\tilde{M}q} \right] = \frac{K}{\tilde{M}} = \frac{1}{\bar{M}}$$

$\theta^* \sim g$, and consider (Z, θ^*)

$$Z|\theta^* \sim \text{Bern} \left(\frac{\tilde{P}(\theta^*|y)}{\tilde{M}g(\theta^*)} \right)$$

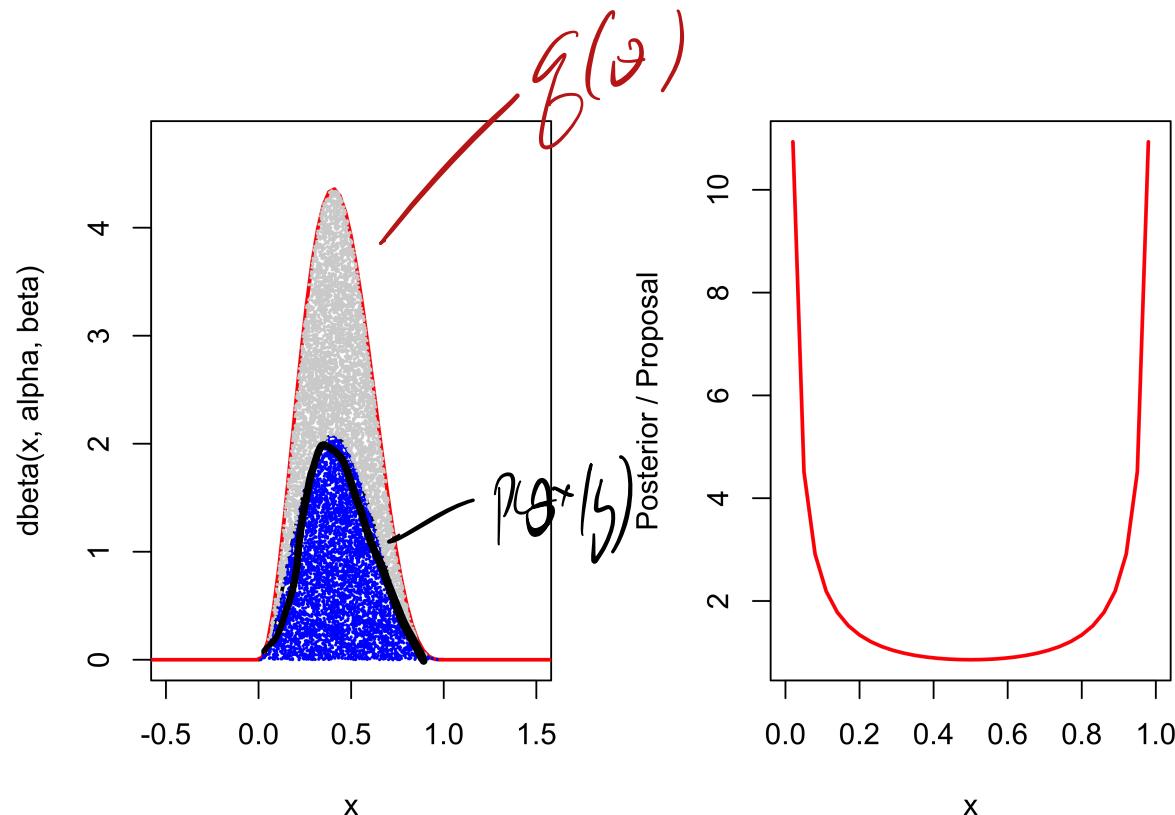
$$P(Z, \theta) = q(\theta) \left(\frac{\tilde{P}(\theta|y)}{\tilde{m}q(\theta)} \right)^Z \left(1 - \frac{\tilde{P}(\theta|y)}{\tilde{m}q(\theta)} \right)^{1-Z}$$

$$P(\theta | Z=1) = \frac{P(Z=1, \theta)}{P(Z=1)} = \frac{q(\theta) \frac{\tilde{P}(\theta|y)}{\tilde{m}q(\theta)}}{K/\tilde{m}}$$

$$= \frac{\tilde{P}(\theta|y)}{K} = p(\theta|y)$$

Proposal most envelope target

Sometimes its not obvious...



Target density in black, proposal density in red

The German Tank Problem



The German Tank Problem

- US and allies wanted to estimate how many tanks were produced by the Germans during WWII
- They used the serial numbers on captured or destroyed tanks to estimate tank production
- Assuming serial numbers, Z_i , are ordered from 1 to n , where n is the number of tanks produced
- How do we estimate n ?

The German Tank Problem

- Assume we observe the serial number of only one tank, Z
- What is a plausible sampling model? How do we write the likelihood?
- What is a reasonable prior?

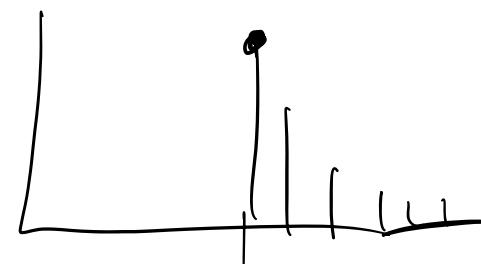
Observe $y \sim \text{Discrete Unif}(1, 2, \dots, n)$

The German Tank Problem

- Assume Z is uniformly distributed between 1 and n
- Specify a Poisson prior distribution
 - Assume a priori a rate of 200 tanks / month
- Use Monte Carlo (rejection sampler) to summarize the results. What is a reasonable proposal?

MLE for n ? $L(n) = \frac{1}{n} I[n > y]$

$$\hat{n}_{MLE} = 6$$
$$\hat{n} = 2$$



$$n \sim \text{Pois}(\lambda)$$

$$P(n|y) \propto \frac{1}{n!} I[n \geq y] \frac{e^{-\lambda} \lambda^n}{n!}$$

Propose from $g \sim \text{Pois}(\lambda)$

$$\tilde{\lambda} = \max_n \frac{P(n|z)}{g(n)} = \max_n \frac{\frac{1}{n!} e^{-\lambda} \lambda^n}{\frac{e^{-\lambda} \lambda^n}{n!} I[\cdot]}$$

$$= \max_n \frac{1}{n} I[n \geq y]$$

$$= \frac{1}{y}$$

$$\tilde{P}(n|y) = \frac{\frac{1}{n} e^{-\frac{1}{n}} \frac{1}{n!} I[n]}{e^{-\frac{1}{n}} \frac{1}{n!}} y.$$

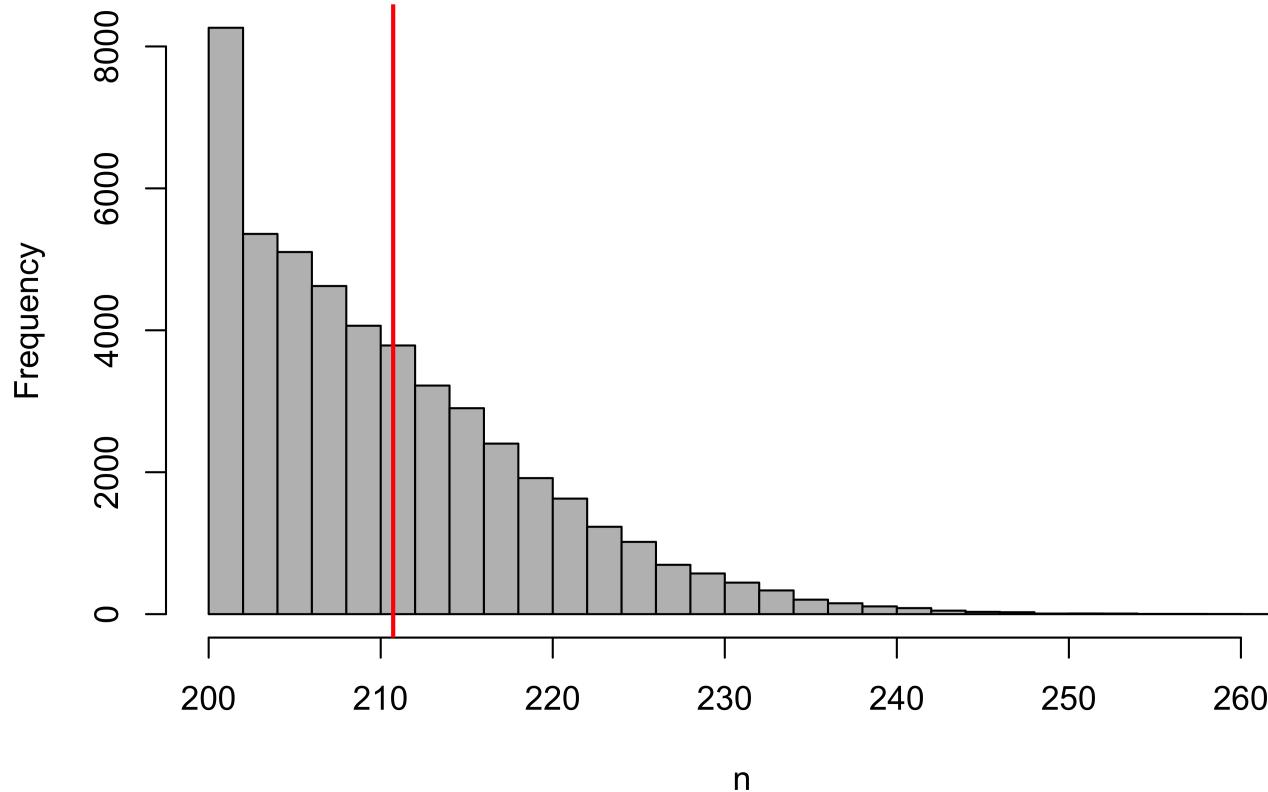
$$= \sum_n I[n \leq y]$$

mom y

The German Tank Problem, $\hat{Z} = 200$

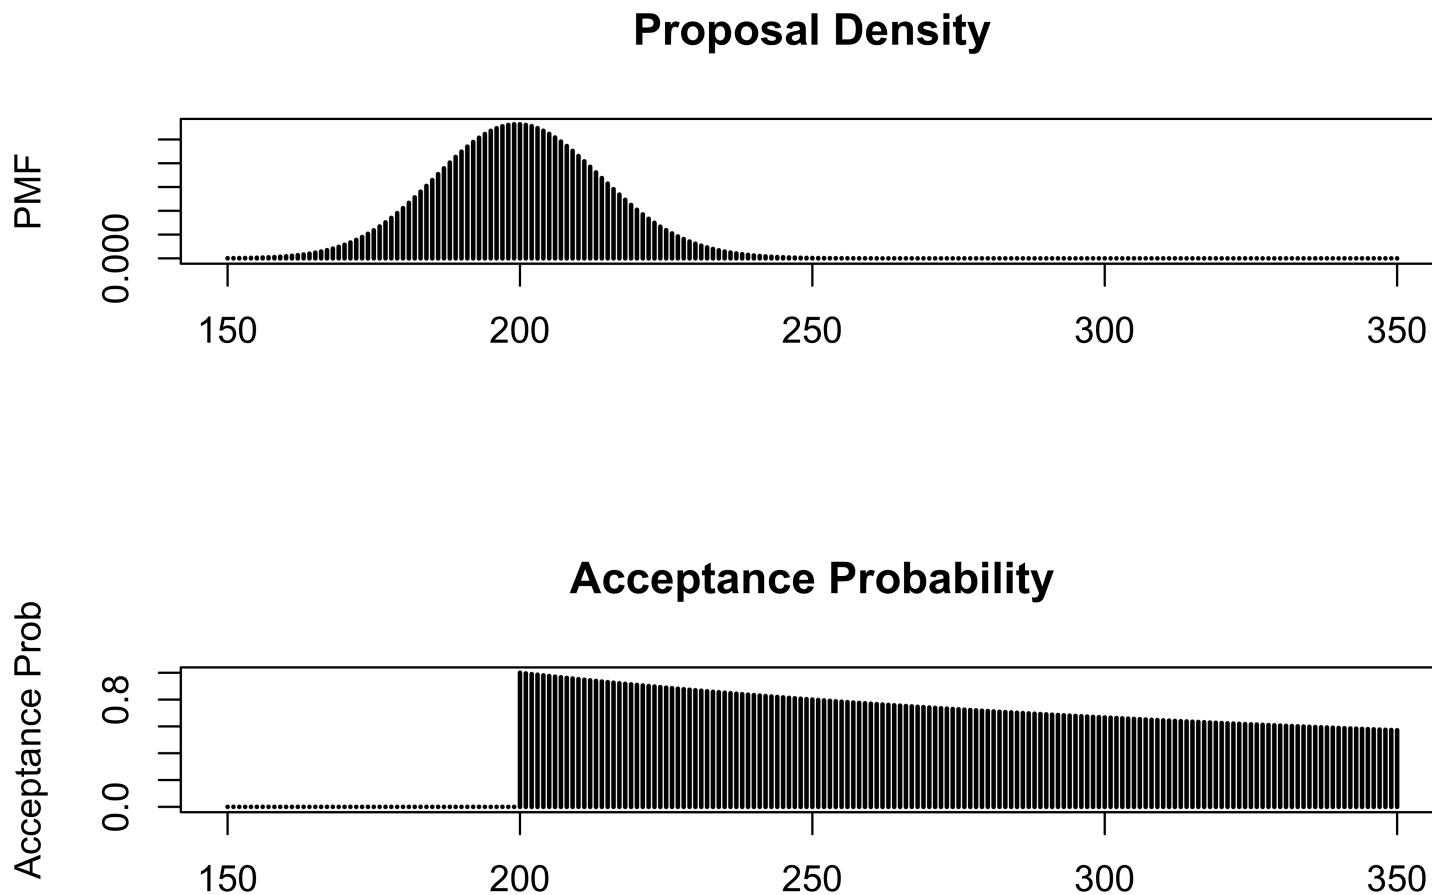
Handwritten note: $\hat{Z} = 200$

Histogram of posterior samples (posterior mean is 211)

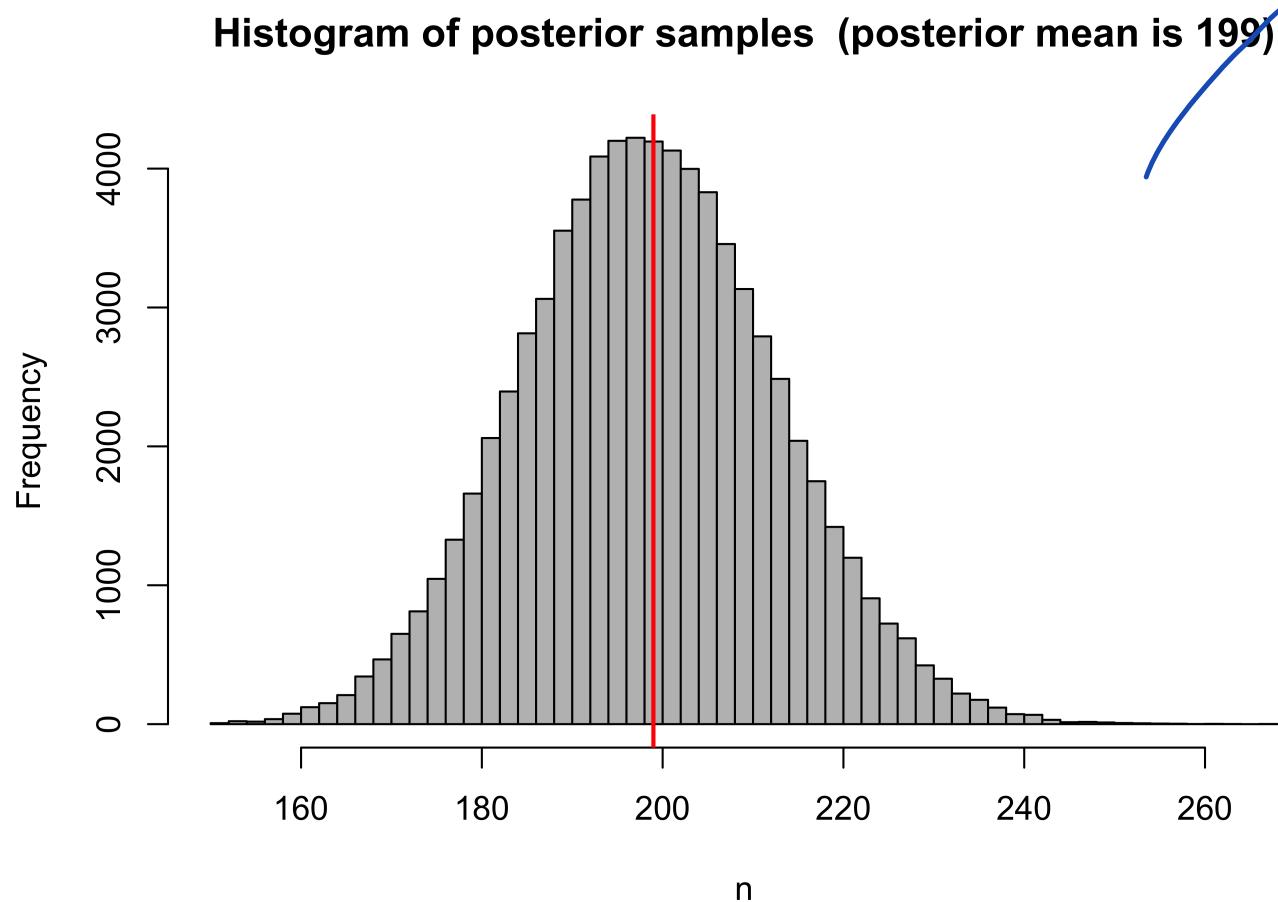


Acceptance rate: 0.48257

The German Tank Problem, $Z = 200$



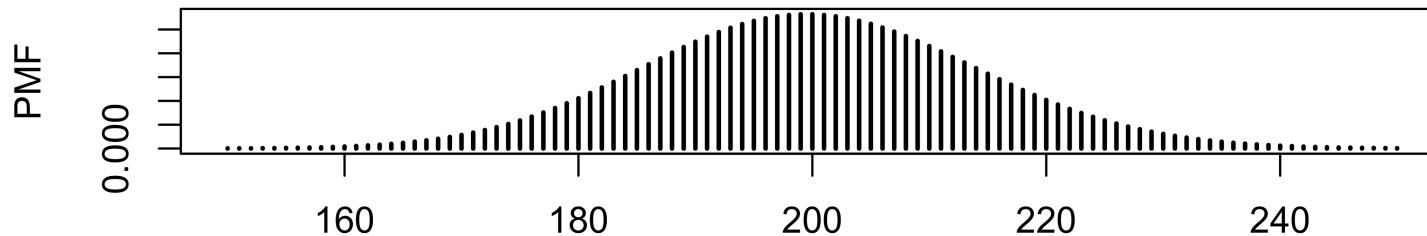
The German Tank Problem, $Z = 150$



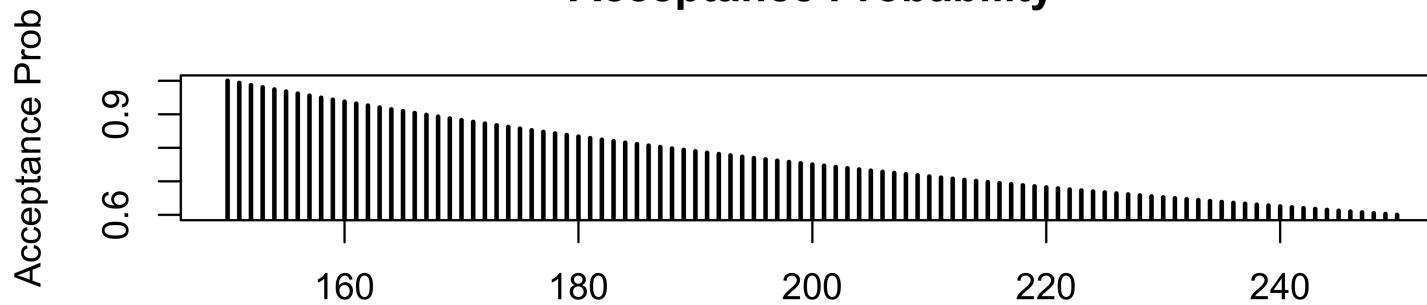
Acceptance rate: 0.75292

German Tank Problem, Z=150

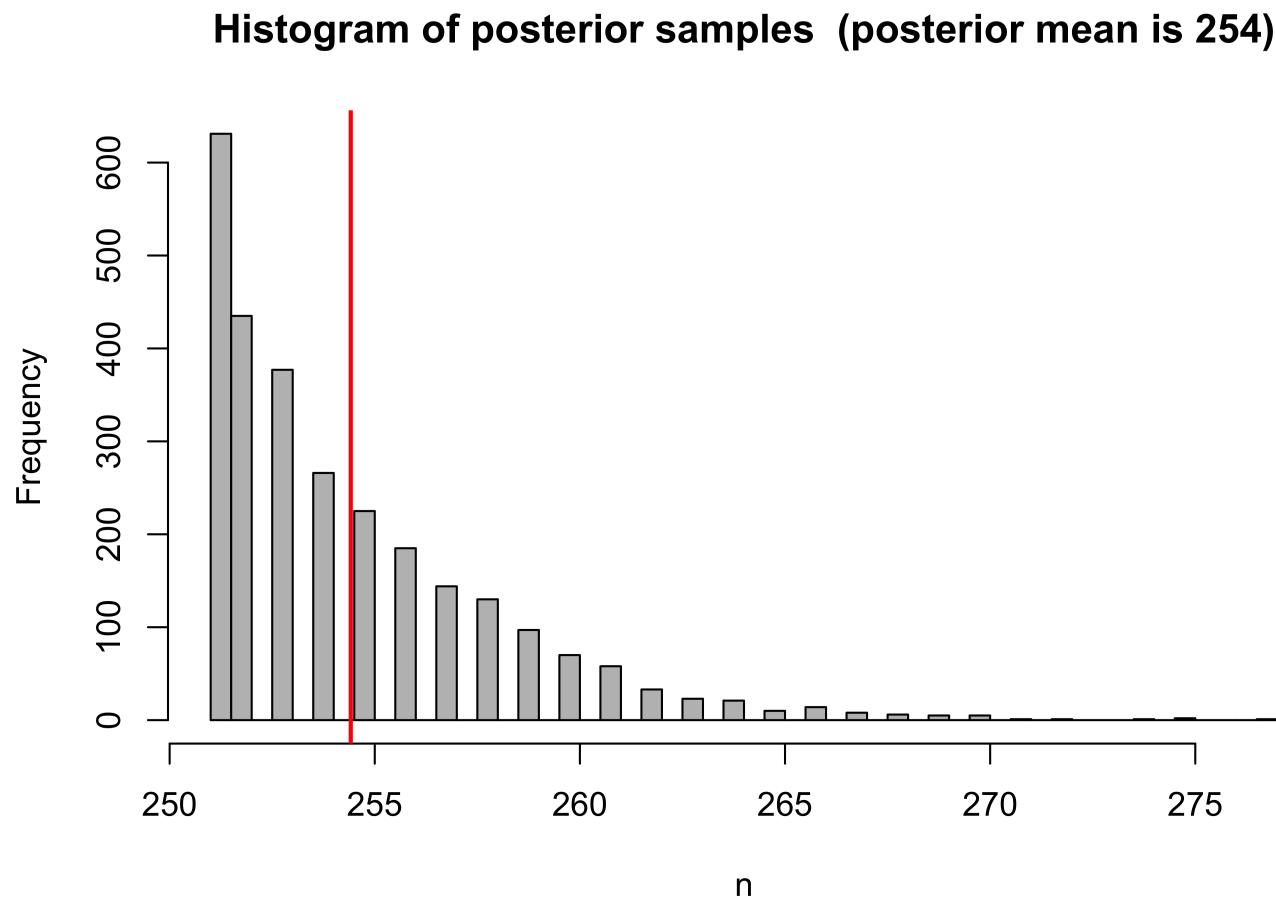
Proposal Density



Acceptance Probability



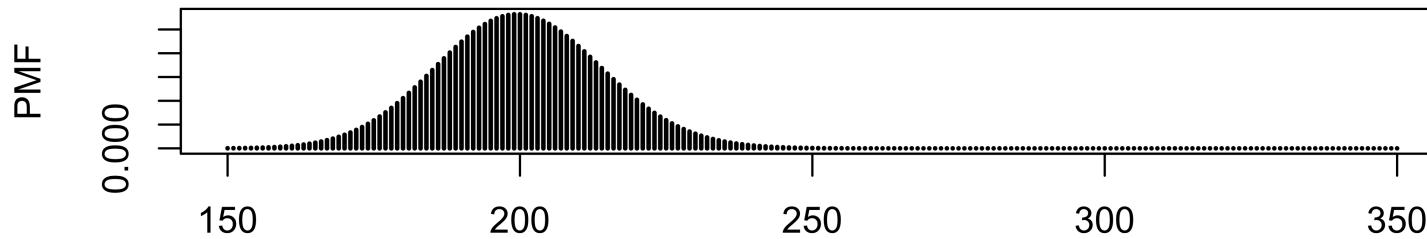
The German Tank Problem, Z = 250



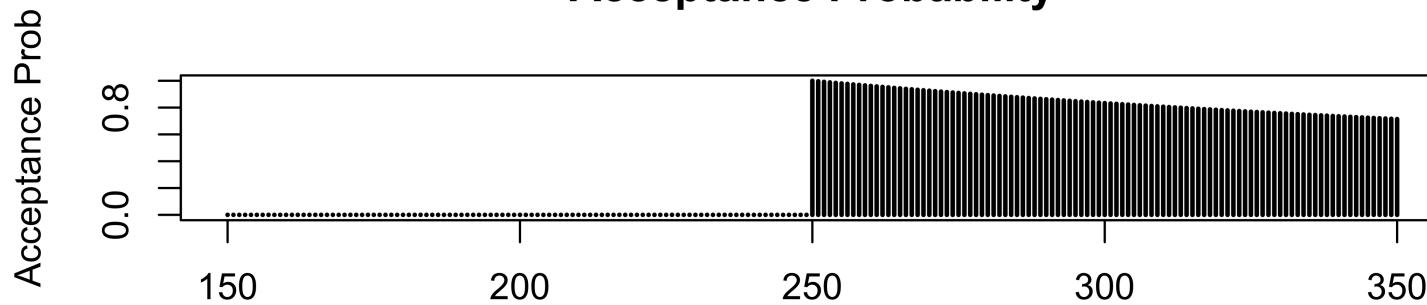
Acceptance rate: 2.749^{-4}

The German Tank Problem, $Z = 250$

Proposal Density



Acceptance Probability



Importance Sampling

- With rejection sampling we get actual samples from the target density
- Downside: often throw away many samples. Inefficient if rejection rate is high.
 - Main challenge is finding a valid proposal density and one that minimizes the number of rejections
- Importance sampling: use all samples from proposal to compute expected values.
- Weight the *importance* of each sample by how representative it is of the target distribution

$$\text{Want} + E_{P(\theta|y)}[h(\theta)] \stackrel{\text{LRTVS}}{=} \dots$$

$$= \underbrace{\int h(\theta) \frac{P(\theta|y)}{q(\theta)} q(\theta) d\theta}_{P/q = w(\theta)}$$

$$= E_{q(\theta)} \left[\frac{h(\theta) P(\theta)}{q(\theta)} \right] \quad P/q = w(\theta)$$

$$\approx \frac{1}{S} \sum_s w(\theta_s) h(\theta_s) \quad \theta_s \stackrel{\text{iid}}{\sim} q(\theta)$$

$$E_{q(\theta)}[w(\theta)] = \int \frac{P(\theta|y)}{q(\theta)} q(\theta) d\theta = 1$$

$$\tilde{w}(\theta) = \frac{\tilde{P}(\theta|y)}{q(\theta)} / \sum \frac{\tilde{P}(\theta|y)}{q(\theta)}$$

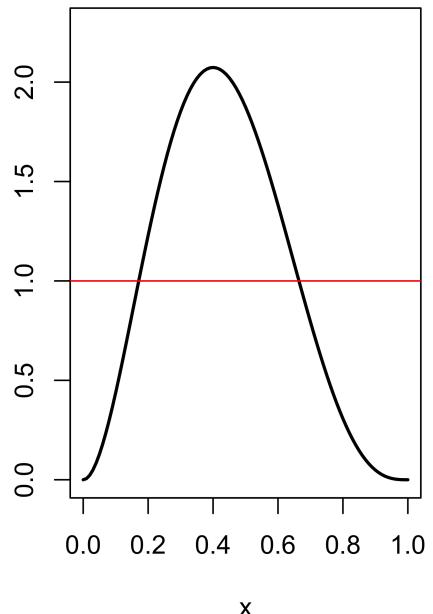
Importance Sampling

- Examine histogram of log importance weights
- Want the largest weights to be not too large
- Effective sample size: $\frac{1}{\sum_{s=1}^S \tilde{w}(\theta^s)^2}$ where $\tilde{w}(\theta^s)$ are the normalized weights

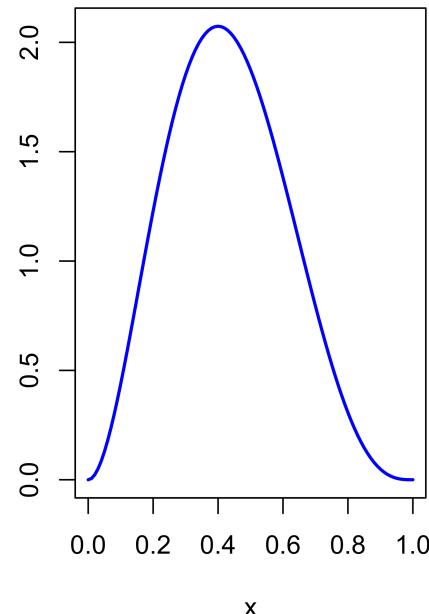
Importance Sampling

- Target: a $\text{Beta}(3, 4)$
- Proposal: Uniform

Target and Proposal Densities

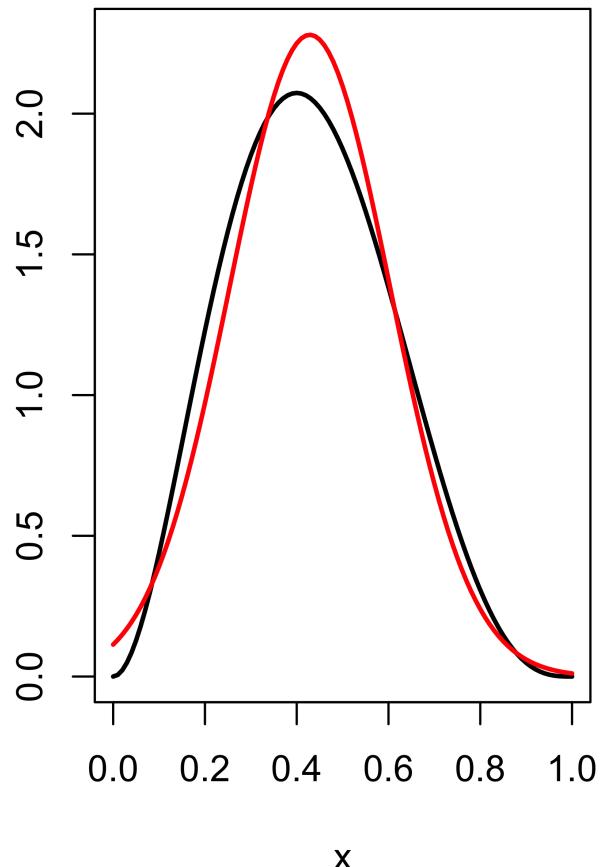


Importance Weights

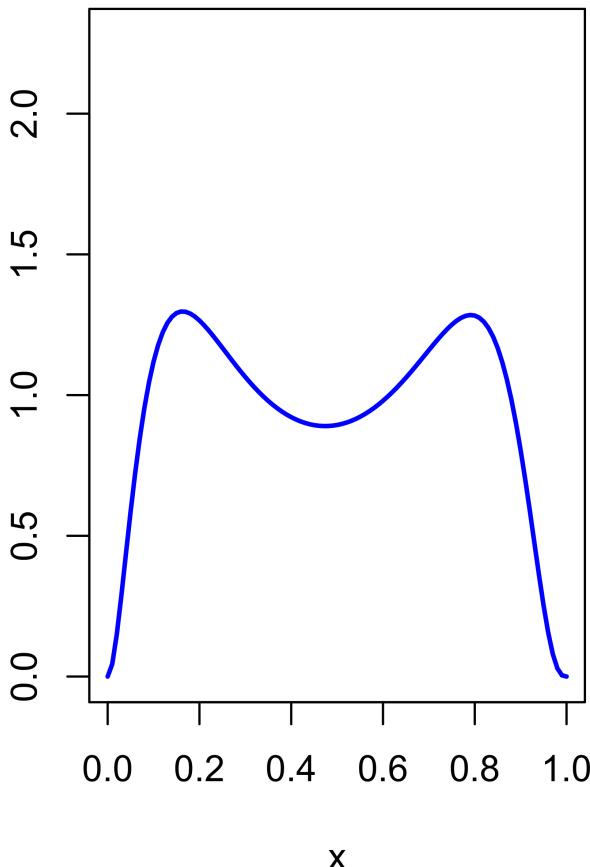


Importance Sampling

Target and Proposal Densities

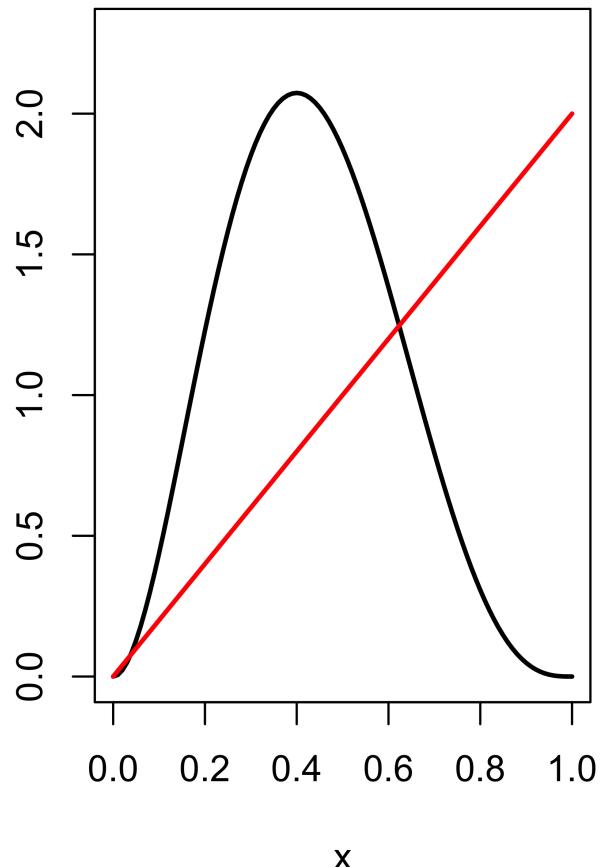


Importance Weights

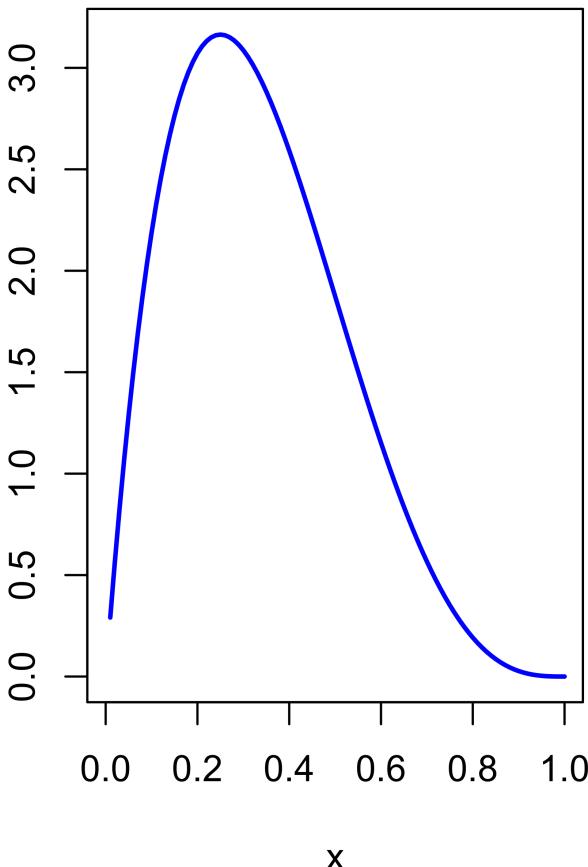


Importance Sampling

Target and Proposal Densities

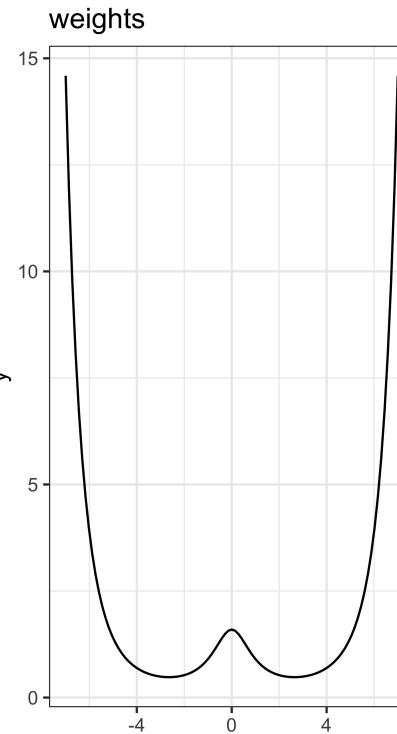
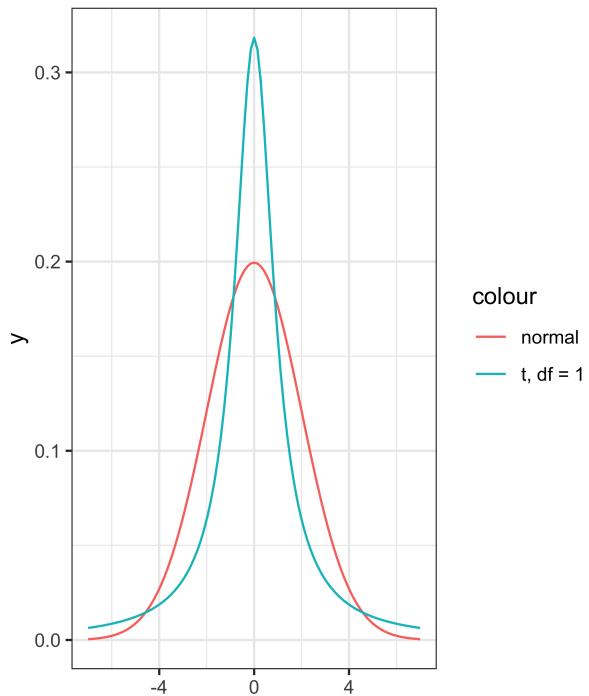


Importance Weights



Importance Sampling

- Target: Cauchy(0, 1)
- Proposal: Normal(0, 3)



Importance Sampling

```
1 df <- 1
2
3 proposal_draws <- rnorm(1e5, sd=2)
4 proposal_draws <- sort(proposal_draws)
5
6 w <- dt(proposal_draws, df=df) / dnorm(proposal_draws, sd=2)
7 w <- w/sum(w)
8
9 q <- 0.25
10 idx <- match(FALSE, cumsum(w) < q)
11 qt(q, df=df)
```

```
[1] -1
```

```
1 proposal_draws[idx]
```

```
[1] -0.860783
```

```
1 q <- 0.05
2 idx <- match(FALSE, cumsum(w) < q)
3 qt(q, df=df)
```

```
[1] -6.313752
```

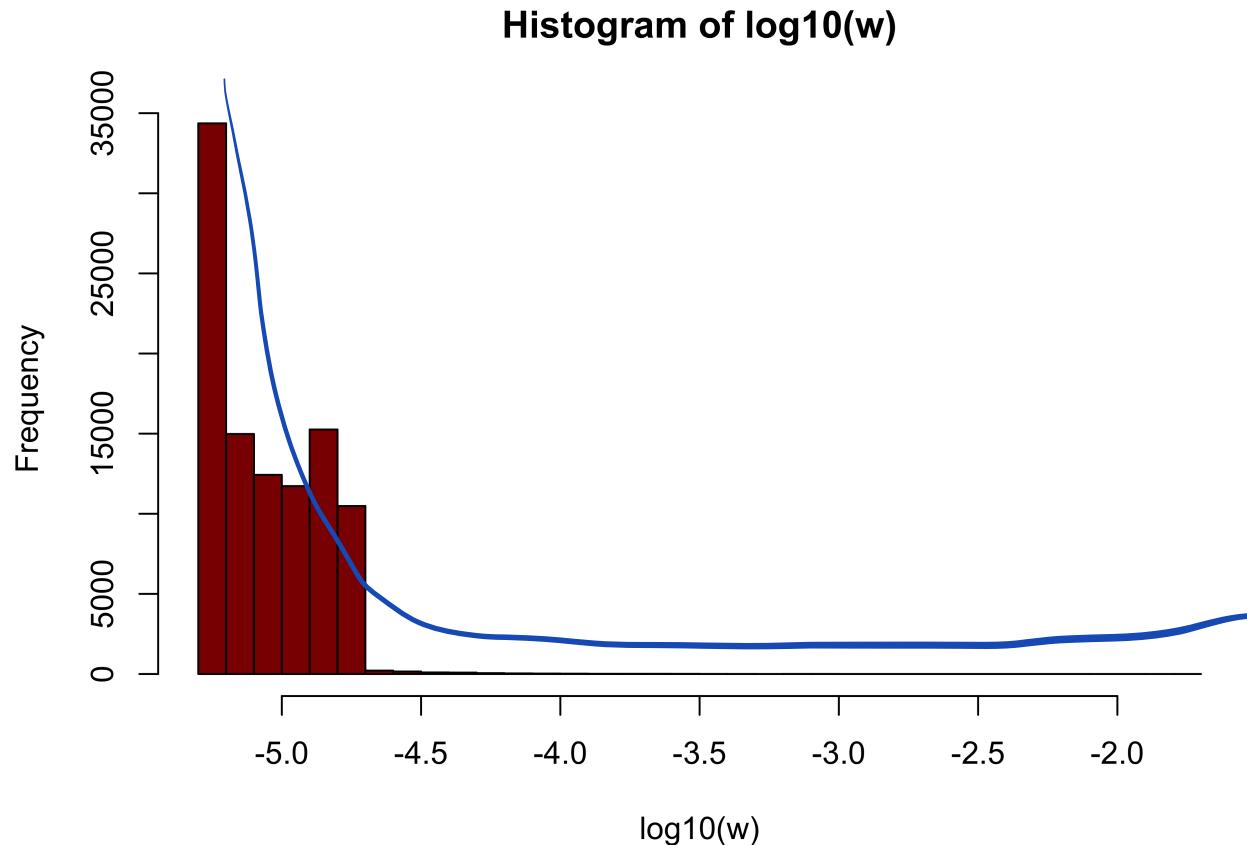
```
1 proposal_draws[idx]
```

```
[1] -3.738269
```

Importance Sampling

```
[1] "99% quantile of log10(w) is -4.773680"
```

```
[1] "max of log10(w) is -1.769665"
```



A Gamma Puzzle

- Goal: report the exact mean of a $\text{Gamma}(a, b)$ using only one random draw
- Allowed to generate a random draw from a proposal distribution of your choosing
- You can evaluate the gamma density and the proposal density at the generated value

$$\int x p(x) dx = \int_0^a x^b \frac{e^{-bx} x^{a-1}}{\Gamma(a)} dx$$

Propose $\text{Gam}(a+1, b)$

$$= \int x^b e^{-bx} x^{a+1} \frac{b^{a+1} e^{-bx} x^a}{\Gamma(a+1)} dx$$
$$\underbrace{b^{a+1} e^{-bx} x^a}_{\Gamma(a+1)}$$
$$= \int_0^{N \times} b \frac{e^{-bx}}{\Gamma(a+1)} x^a dx$$

Sketch

A Gamma Puzzle

```
1 ## Generate 10 draws from a gamma 5, 2
2 proposal_draws <- rgamma(10, shape = 5, rate=2)
3
4 ## Target a gamma 4, 2
5 w <-
6   dgamma(proposal_draws, shape=4, rate=2) /
7   dgamma(proposal_draws, shape=5, rate=2)
8
9 w*proposal_draws
```

[1] 2 2 2 2 2 2 2 2 2 2